


Building a high-retention Spotify video ▾



im working on a video i want to create visuals (actual high quality visuals that tell the story of what i'm talking about) the video is about spotify and its 1 hour + i want you to build the video, using tools at your disposal, i want to be able to preview the video too here so you might create a previewer i can use to play it but the goal right now is to figure out how we can actually get this shit moving cause rn i'm thinking of the best approach and style to make this video there are certain things i have in mind: 1. increase retention 2. scenes change every 3 seconds maximum of 10if context needs 3. match the script so now you gonna give me suggestion based on the transcript (i'll attach the audio when we finalize)

Ran 7 commands, viewed a file

Now I have a very solid picture of your video. Let me give you my full strategic breakdown.

Now I have everything I need. Let me build you a comprehensive strategy widget and video structure.

Video length	SRT segments	Word-level timestamps	Scene changes (~3s)
68 minutes	1,064	10,481	~1,360

STORY ARC — 7 CHAPTERS

0:00 – 1:33 The hook You pay more, artists get less, Spotify keeps the diff. Cold open, no intro. High energy	1:33 – ~12 min The price hike receipts 4 hikes in 4 years. \$9.99 → \$14.99. Fake feature bundling. Psychological tricks. Data heavy	~12 – 28 min Where the money actually goes \$10B lie. Per-stream math. Label carve-outs. Discovery Mode payola. Data heavy	~28 – 42 min Artists disappearing Middle class of music gone. 1K-stream zero-out. Fraud incentive. The Chance era is over. Narrative
~42 – 55 min The failed pivots \$1B podcast bet. Killed shows. No lossless in 2026. AI noise flooding catalog. Analysis	~55 – 65 min Daniel Ek & the stock game CEO selling shares. Defense company pivot. The platform quietly decays while he exits. Expose	~65 – 68 min The verdict 3 numbers, 3 trajectories. What you can actually do. Buy Bandcamp. End. CTA	

VISUAL SCENE TYPES — MAPPED TO SCRIPT

📊 Animated stats	Counters that tick up on screen. "\$9.99 → \$10.99 → \$11.99 → \$12.99 → \$14.99" with year stamps. Per-stream rate dropping from 0.7¢ to 0.3¢. Revenue vs artist payout diverging lines.
📄 Receipt / doc mockups	Fake Spotify invoice revealing where \$11.99 goes. Legal docs. Blog post screenshots. Label contract breakdown. All styled like receipts — monospace, paper texture.
▶ App UI recreations	Spotify UI mockup with Discovery Mode toggle + "lower royalty" warning. Playlist with major label vs indie payouts side-by-side. "1,000 streams = \$0.00" payout screen.
📰 Headlines / press clippings	Animated newspaper-style cards. "Spotify hikes prices again." "Spotify cancels exclusive shows." "Ek invests in defense." Stack in chronologically, each with a date stamp.
📈 Flow diagrams	\$1 → Spotify (30¢ off top) → labels → distributor → publisher → songwriter → artist (trace of a cent). Revenue share pool animation showing major labels eating 70% each cycle.
👤 Human-scale math	Artist avatar: 2,000 listeners. 60,000 streams. \$180/mo. Before fees. Visual subtraction until it hits near zero. Side-by-side with Spotify's profit in same frame.
🌐 Geography / rate maps	World map showing per-stream value by country. Same song, same stream — 0.5¢ in Norway, 0.07¢ in Nigeria. Animated fill as narration covers each country.

RETENTION MECHANICS BUILT INTO THE EDIT



3-second max cuts

Every visual holds max 3s before a cut or motion event. Word-level JSON means we can sync cuts to syllables, not just sentences.



Pattern interrupts

Every 90 seconds: new visual format (stat → mockup → diagram → headline). Keeps brain from habituating to the rhythm.



Text on screen

Key numbers and phrases always on screen when spoken. "50% more" appears huge. "Artists make \$180/mo" shown as a receipt. Don't just say it — show it.



Escalating tension

Each chapter reveals something worse than the last. Price hikes → payout collapse → middle class gone → AI noise → CEO exits. Payoff at the end.



Chapter anchors

Short "chapter card" between sections (1–2 seconds). Gives viewers who scrub back a landmark. Also signals "there's more coming" to keep them watching.



The receipt motif

Script says "read the receipt" — lean into this as a recurring visual motif. Every data reveal styled like a line item on a receipt. Brand cohesion + mnemonic hook.

RECOMMENDED VISUAL STYLE

Color palette

Black + white base. Spotify green (#1DB954) used sparingly — only when being critiqued. Red used for "bad" numbers. Monochrome default so green pops when it appears.

Typography

Big, bold, condensed sans for stats. Monospace for receipts and docs. No decorative fonts. Text should feel like audit paperwork, not a marketing deck.

Motion style

Fast cuts (not fades). Slide-in from direction that matches context (price going up = slides from bottom). Numbers count up. Charts draw left-to-right. No spinning logos.

Energy pacing

Hook is the fastest (cuts every 1–2s). Data chapters breathe a bit more (2–3s). Verdict is slow and deliberate. Match voiceover rhythm precisely using word timestamps.

[Build the video previewer ↗](#)

[Scene breakdown: first 5 min ↗](#)

[Design price hike scenes ↗](#)

Here's the full strategic picture based on your transcript.

What you're working with is exceptional source material. The script is already structured like a great video essay — it has a clear cold open, escalating chapters, a villain (Daniel Ek), and a punchy verdict. The word-by-word JSON timestamps mean we can sync visuals to *syllables*, not just sentences — that's a huge advantage for the kind of tight, high-energy editing you want.

The biggest stylistic decision you need to make before we build anything is whether this video is:

1. **Documentary/analytical** — clean, minimal, black-and-white data viz, feels like an audit (think Johnny Harris or Wendover). The "receipt" motif in your script is perfect for this.

2. **Confrontational/punchy** — aggressive cuts, big text, more YouTube-essay energy (think Coffeozilla or Internet Historian). Leans into the scam angle hard.
3. **Hybrid** — starts cold and analytical, gets more heated as the Daniel Ek section lands.

My honest take: **option 3 is the move**. Your script earns the right to get angry by being rigorous first.

The previewer is fully buildable — I can make an interactive scrubber in this chat where you paste a timestamp and see the designed scene. The word-level JSON makes this precise. Ready to build it the moment you confirm direction on style.

What do you want to dig into first?


So basically, what I'm trying to achieve is something that would ensure there's retention. So I was thinking of doing this doodle style animation, black and white. And, um, like a pencil drawn, not necessarily black and white. Sorry. If you feel undrawn, colorful, but not too colorful. Like, Spotify, local color can be there. and, um, other colours and elements can be there. Like, they can be colour. You should feel and drawn and should be drawn in real time. And I want to do this with GLM five point two. So I was thinking, how would I go about this? because I want I want there to be element drawn into, um, the canvas. Like, as I'm talking, the elements are being drawn. So I'm talking the elements are being drawn for retention. I think that is the best strategy, and they should not feel like a slideshow. It should feel like an actual to do animation, but I'm seeing something to be using the AI agent. I don't really know how to go about it.

This is a great vision. Let me break down exactly how this works before we build anything.


The core idea: As your voiceover plays, an AI agent (GLM or similar) is given the current transcript segment and draws elements onto a canvas in real-time — lines appearing, text being written, characters sketched — all synced to your word timestamps.

Let me show you the architecture first, then we build.

your word_by_word.json already has everything needed to drive this



what the AI agent sees and returns — per 3-second window



Here's the real breakdown of how this works technically, and the honest truth about GLM.

The pipeline in plain terms:

Your `word_by_word.json` already has every word with a start and end timestamp. The playback engine reads the audio, and every ~3 seconds it fires a "draw this" event to an AI agent. The agent gets told what's being said right now and returns a sequence of strokes — paths, text, shapes — each with a time offset so they appear as the words are spoken. The HTML Canvas API draws them stroke by stroke, not all at once, creating the "being drawn in real time" effect.

About GLM 5.2 specifically — this is a smart call but needs clarification. GLM (Zhipu AI's model) is strong at instruction-following and can return structured JSON. What you'd use it for is generating the *description* of what to draw and the stroke coordinates in JSON format. It won't directly "draw" — it tells your canvas what to draw. The challenge is that GLM needs to return precise coordinate paths, which requires either a visual model that understands canvas coordinates, or you pre-define a set of drawing

primitives (Spotify logo, price tag, arrow, person sketch, bar chart) and let GLM *compose* them rather than invent new strokes from scratch.

The two real approaches to choose between:

1. **AI-guided composition** (simpler, more reliable) — You build a library of ~50 hand-drawn SVG primitives (Spotify circle, coins, arrows, person, headphones, bar chart, price tag, etc.). GLM reads the transcript and picks which primitives to place and animate. You control the hand-drawn style, GLM controls the storytelling logic.
2. **AI-generated strokes** (ambitious, harder) — GLM or a vision model generates actual SVG path data that gets drawn progressively. This is harder to keep consistent but gives full creative freedom. You'd need fine-tuning or very tight prompting to keep the doodle style coherent.

My recommendation for your video: Start with approach 1. Build the previewer here first so you can see what each 3-second window should look like, then we build the actual drawing pipeline. Want me to build the interactive previewer now — you click any point in the 68-minute transcript and see a mockup of what the doodle scene would look like?

How do I build the playback engine that reads timestamps and fires events?

Here's exactly how to build it.

LIVE DEMO — PASTE YOUR WORD_BY_WORD.JSON OR USE SAMPLE DATA

II Pause

Reset

Load sample

0.07s

You've

You've paid Spotify more money every single year since

Event log

money 2.07s word every 2.30s word single 0s draw_window 0 → AI: "You've paid Spotify more money every single year" 0.03s word You've 0.03s segment You've paid Spotify more money every sin... 0.19s word paid 0.59s word Spotify 1.30s word more 1.54s word money 2.04s word every 2.36s word single 2.84s word year 3.51s word since 3.95s word 2023. 4.30s word And 4.70s word artists 5.10s word are 5.50s word getting 3.95s segment 2023. And artists are getting paid less... 3s draw_window 1 → AI: "since 2023. And artists are getting" 6.00s word paid 6s draw_window 2 → AI: "paid less per stream than ever." 0s draw window 0 → AI: "You've paid Spotify more money every single

THE CORE ENGINE — COPY THIS

```
class PlaybackEngine { constructor(words, segments) { this.words = words; // word_by_word.json
array this.segments = segments; // SRT segments array this.listeners = {}; this.startedAt = null;
this.offsetMs = 0; this.wordIdx = 0; this.segIdx = 0; this.drawIdx = 0; this.raf = null;
this.playing = false; } // Subscribe to events on(event, cb) { (this.listeners[event] ??=
[]).push(cb); return this; } emit(event, data) { (this.listeners[event] ?? []).forEach(cb =>
cb(data)); } play(audioEl) { this.audio = audioEl; this.playing = true; this.tick(); } tick() { if
(!this.playing) return; const now = this.audio.currentTime; // seconds // Fire word events while
(this.wordIdx < this.words.length && this.words[this.wordIdx].start <= now) { this.emit('word',
this.words[this.wordIdx]); this.wordIdx++; } // Fire segment events (SRT lines) while (this.segIdx
< this.segments.length && this.segments[this.segIdx].start <= now) { this.emit('segment',
this.segments[this.segIdx]); this.segIdx++; } // Fire draw events every 3 seconds const drawWindow
= Math.floor(now / 3); if (drawWindow > this.drawIdx) { this.drawIdx = drawWindow;
this.emit('draw_window', { windowIndex: drawWindow, timeStart: drawWindow * 3, timeEnd: drawWindow
* 3 + 3, words: this.words.filter( w => w.start >= drawWindow*3 && w.start < drawWindow*3+3 ) });
} this.raf = requestAnimationFrame(() => this.tick()); } pause() { this.playing = false;
cancelAnimationFrame(this.raf); } seek(seconds) { this.wordIdx = this.words.findIndex(w => w.start
>= seconds); this.segIdx = this.segments.findIndex(s => s.start >= seconds); this.drawIdx =
Math.floor(seconds / 3); } }
```

WIRE IT UP — 10 LINES

```
const engine = new PlaybackEngine(words, segments); const audio = document.getElementById('my-
audio'); // Highlight word on screen engine.on('word', ({text, start, end}) => {
document.getElementById('cur-word').textContent = text; }); // Show current subtitle line
engine.on('segment', ({text}) => { document.getElementById('cur-seg').textContent = text; }); //
🔑 This is where you call the AI agent engine.on('draw_window', async ({words, timeStart}) => {
const phrase = words.map(w => w.text).join(' '); const cmds = await askAIToDraw(phrase,
```

```
timeStart); canvas.renderCommands(cmds); // stroke by stroke }); engine.play(audio);
```

SEEK TO ANY TIMESTAMP

```
// Jump to 22 seconds (the "Spotify raised prices" line) audio.currentTime = 22; engine.seek(22);
// Or sync to audio scrubbing automatically audio.addEventListener('seeked', () => {
  engine.seek(audio.currentTime); });
```

Hit "Load sample" then "Play" to see all three event types firing in real time.

The three events and what you do with each:

The `word` event fires for every word as it's spoken — use this to highlight text on screen if you want karaoke-style subtitles. The `segment` event fires per SRT line — useful for showing the full subtitle. The `draw_window` event is the critical one — it fires every 3 seconds and tells the AI agent what phrase is being spoken right now, so it can generate the doodle for that window *just before* it needs to appear.

One key thing about timing: In a real implementation you want to fire `draw_window` about 1–2 seconds *early* — before the window actually starts — so the AI has time to respond and the first strokes begin appearing right on cue. Change `Math.floor(now / 3)` to `Math.floor((now + 1.5) / 3)` and you get a 1.5-second lookahead buffer.

To connect your actual files, you load the JSON and parse the SRT once on startup, then hand them to the engine. The audio element is just a standard `<audio>` tag — the engine reads `audio.currentTime` on every animation frame, which gives you sub-frame accuracy without any timers.

Want me to build the next piece — the canvas renderer that takes draw commands and animates strokes appearing one by one in the doodle style?


yes

LIVE DOODLE CANVAS — PICK A SCENE FROM YOUR TRANSCRIPT

 Price hikes

 Artist payout

 Per-stream math

 Polite scam


done — hit replay or pick another scene

 Replay

Clear

Speed: 1×

THE RENDERER CLASS — DROP THIS INTO YOUR PROJECT

```
class DoodleRenderer { constructor(canvas) { this.ctx = canvas.getContext('2d'); this.queue = [];
this.busy = false; this.speed = 1; // multiply all durations by 1/speed this._applyPaperTexture(); } //
Main entry — pass AI-returned command array renderCommands(cmds) { this.queue.push(...cmds); if
(!this.busy) this._next(); } async _next() { if (!this.queue.length) { this.busy = false; return; }
this.busy = true; const cmd = this.queue.shift(); await this._execute(cmd); this._next(); } async
_execute(cmd) { switch (cmd.type) { case 'stroke': return await this._drawStroke(cmd); case 'text':
return await this._drawText(cmd); case 'circle': return await this._drawCircle(cmd); case 'wait':
return await this._wait(cmd.ms); case 'clear': return this._clear(); } } // Animate a path point-by-
point (hand-drawn feel) async _drawStroke({points, color='#222', width=2.5, duration=600}) { const ctx
= this.ctx; const ms = duration / this.speed; const delay = ms / points.length; ctx.strokeStyle =
color; ctx.lineWidth = width; ctx.lineCap = 'round'; ctx.lineJoin = 'round'; for (let i = 1; i <
points.length; i++) { ctx.beginPath(); ctx.moveTo(points[i-1].x, points[i-1].y);
ctx.lineTo(points[i].x, points[i].y); ctx.stroke(); await this._wait(delay); } } // Typewriter-style
hand font text async _drawText({text, x, y, size=28, color='#222', duration=500}) { const ctx =
this.ctx; const delay = (duration / this.speed) / text.length; ctx.fillStyle = color; ctx.font = `700
${size}px 'Caveat', cursive`; let drawn = ''; for (const ch of text) { drawn += ch; ctx.clearRect(x -
2, y - size, ctx.measureText(drawn).width + 40, size + 8); ctx.fillText(drawn, x, y); await
this._wait(delay); } } _wait(ms) { return new Promise(r => setTimeout(r, ms)); } }
```

WHAT THE AI AGENT RETURNS — COMMAND FORMAT

```
// For the line: "Spotify raised its prices in 2023" // AI returns this array of draw commands: [ {
type: 'wait', ms: 200 }, // Draw Spotify circle logo { type: 'circle', cx: 120, cy: 180, r: 60, color:
'#1DB954', fill: true, duration: 400 }, // Write the price going up { type: 'text', text: '$9.99', x:
240, y: 160, color: '#222', size: 36, duration: 300 }, { type: 'text', text: '→', x: 340, y: 160,
color: '#E34948', size: 36, duration: 150 }, { type: 'text', text: '$10.99', x: 370, y: 160, color:
'#E34948', size: 36, duration: 300 }, // Underline stroke { type: 'stroke', color: '#E34948', width: 3,
duration: 400, points: [{x:240,y:170},{x:245,y:172},...,{x:480,y:170}] }, ]
```

Click any of the 4 scenes — they're all pulled directly from your transcript. Hit the speed toggle to see it at 2× or 4× while building.

What you just saw is the complete renderer. Three things make the doodle feel hand-drawn rather than like a slideshow:

The `Caveat` font (Google Fonts, free) is a real handwriting typeface — it's what makes the text look sketched rather than typed. The stroke points have tiny random jitter baked in (`±1.5px`) so lines are never perfectly straight. And everything draws sequentially — nothing appears all at once.

The command format is the key interface between your AI agent and the canvas.

When GLM or Claude gets the current phrase, it returns an array like

`[{type:'circle',...}, {type:'text',...}, {type:'stroke',...}]` — that's it. The renderer just executes them in order. You can add new command types (like `arc`, `arrow`, `scribble`, `shake`) without touching anything else.

What to build next — in order:

First, add the `draw_window` handler from the playback engine to call a real AI endpoint and pass back commands to this renderer. Second, record the canvas output — you can use the `MediaRecorder` API to capture the canvas stream directly as a `.webm` file. Third, build the scene pre-generator that processes your entire 68-minute transcript *ahead of time* and stores all the draw command arrays, so during playback there's zero AI latency.

Want me to build the pre-generator next — the script that reads your full

`transcript.srt`, sends each 3-second window to the AI, and saves a complete `scenes.json` file you can use for instant playback?

you got the logos wrong tho
and rn im not thinking of the renderer i would create that later
just wanna get the ball runing

Fair enough. What's the actual priority right now — getting the AI agent prompt working so it returns the right draw commands for each 3-second window? That's the piece that connects your transcript to the visuals, and once that's solid everything else slots in around it.

Also noted on the Spotify logo — the real one is three horizontal arcs inside a filled green circle, not text. I'll get that right when we build the actual scene library.

So what do you want to lock in first:

1. The AI agent prompt (what you send GLM per window, what you tell it to return)
2. A pre-generator script that processes your full transcript and outputs `scenes.json`
3. Something else entirely

What's on your mind?

Okay. Um, I like what you did. I like what you did, but I'm thinking of linking my guitar press repo. if you can see what the agent is doing also and like be able to change it too because like, um, g l m five is not really as good as planning as you. You literally got exactly what I wanted in terms of how I want the drawing to be. Although the composition can actually do a lot of bettering, the composition can do a lot better, but I like the fact that it's actually undrawn body composition. I want it to actually be visible on

Show more

Got it. Share the GitHub repo link here and I'll fetch it directly and read exactly what GLM built — the renderer, the prompts, everything. Then I can tell you what's actually solid, what needs fixing, and write you the exact prompt GLM needs to produce good draw commands.

On the tech question — I didn't use Rough.js in what I built. Just plain HTML Canvas with:

- `Caveat` font for the handwriting feel
- tiny random jitter on stroke points (`±1.5px`) to make lines wobbly
- drawing point-by-point with small delays instead of all at once

That's genuinely all it takes. Rough.js can work but it adds a dependency and the sketchy effect it produces is a bit mechanical — what I did looks more natural because the randomness is in the *path*, not applied as a filter after the fact.

Drop the repo link and I'll read it now.

<https://github.com/foskigr8/spotifyvideo> this is the token:
github_pat_11BNMXXCQ0bf35ewf9ubCC_5XLyHG3HymonKWYlbOD0Fz
31h0h0tWFbJ2JqM8GxWlu6XXAH6RQnF7qHchu

Ran 7 commands, searched the web

I've read the whole thing. Here's the real picture.

What GLM actually built is genuinely solid. This is not a half-finished prototype — it's a proper motion graphics pipeline. The architecture is clean: `segmenter.ts` converts your SRT into typed scenes, each scene type maps to a React component, and the playback engine is rAF-driven with word-level sync. The `KineticText` component with the Spotify green highlight is essentially the same thing you liked in my demo, just in React/Framer Motion instead of Canvas.

What's actually missing or wrong:

The `chart` and `receipt` scene types are defined in the types but the components don't exist yet — if a scene gets typed as `chart` or `receipt`, it falls through to `KineticText` as the default. That means your price hike charts and the "read the receipt" hero moment are showing as plain text right now.

The `Stage.tsx` and `Transport.tsx` components I haven't read yet — those are the player shell. The `MoneyFlow` animation looks right conceptually but I want to see `Stage.tsx` to understand the canvas sizing (that's likely why things look off on mobile).

The doodle style you want doesn't exist yet at all — what GLM built is a dark documentary aesthetic (black background, Spotify green, film grain). Beautiful, but completely different from the hand-drawn colorful thing you described. These are two separate visual directions.

The key question you need to decide: do you want to keep GLM's dark documentary style as the base and layer the doodle effect on top, or do you want to pivot the whole thing to the doodle canvas approach? They can coexist — documentary for the previewer/editing tool, doodle canvas for the actual exported video.

Want me to read `Stage.tsx` and `Transport.tsx` now, then tell you exactly what to fix and what to build next?

Also — I'd strongly recommend rotating that GitHub PAT now that it's been shared in this conversation. Go to GitHub → Settings → Developer settings → Personal access tokens and regenerate it.

i want you to actually fix it directly and give me prompt to tell glm 5.2 to do cause i didnt even see the work he did in the preview

Let me read the remaining files first then I'll fix everything directly.

Ran 2 commands

I've read everything. Now I know exactly what's broken and what to fix. Let me push the fixes directly.

Ran 2 commands

Good — the data files exist. Now I know the full picture. Let me push all the fixes:

Problems I found:

1. `chart` and `receipt` scene types fall through to `KineticText` — components don't exist
2. The `layout.tsx` almost certainly has no `font-display` / `font-body` defined — so all text falls back to system font, breaking the whole look
3. The `video-range` CSS class on the scrubber is undefined — it renders ugly
4. Mobile layout breaks because `Stage` has no min-height on small screens

Read layout.tsx to check fonts

Found it. `font-display` is never defined — every component using `font-display` (KineticText, StatReveal, Chapter, Split) is silently falling back to system font. That's why it looks broken. Also `font-body` is set to Geist but `font-mono` conflicts. Let me fix everything now.

Ran 4 commands

